Clinical Discontinuity and AI: Restoring Coherence in Healthcare

Boris Kriger[1,*]

[1] *Institute of Integrative and Interdisciplinary Research*

[*]Corresponding author: `boriskriger@interdisciplinary-institute.org`

ORCID: 0009-0001-0034-2903

January 2026

Abstract

Contemporary healthcare systems underperform despite unprecedented medical knowledge accumulation. This paper argues that medicine's fundamental inefficiency stems not from insufficient knowledge but from structural inability to maintain informational coherence across the clinical process. We introduce the *Clinical Continuity Principle*: effective care requires uninterrupted informational integration from symptom presentation through diagnosis, prescription, and treatment execution. Human cognitive limitations—memory constraints, attentional fatigue, bounded processing capacity—render this continuity structurally unattainable at scale. We synthesize empirical evidence demonstrating systematic information loss, responsibility fragmentation, and feedback disruption across clinical pathways. We argue that artificial intelligence, deployed within hybrid human-AI systems, represents a promising mechanism for restoring clinical continuity, while acknowledging significant limitations including algorithmic bias, reliability concerns, implementation barriers, and the risk of creating new forms of cognitive burden. We propose testable hypotheses for empirical validation and position this framework as a conceptual contribution requiring further research. This analysis repositions AI-medicine discourse from technological enhancement to architectural transformation.

Keywords: clinical continuity, healthcare systems, artificial intelligence, cognitive limitations, medical informatics, systems theory, diagnostic accuracy, treatment adherence

1. Introduction

Modern medicine occupies a paradoxical position. Never before has humanity possessed such extensive knowledge about disease mechanisms, diagnostic markers, therapeutic interventions, and preventive strategies. A 2011 projection estimated that medical knowledge would double approximately every 73 days by 2020 (Densen, 2011)—an illustrative indicator of exponential growth rather than a precise current figure, but one that captures the overwhelming scale of information accumulation. Clinical trials generate increasingly granular efficacy data, and technological innovations continuously expand diagnostic capabilities. Yet healthcare systems worldwide consistently fail to translate this knowledge into proportionally improved outcomes.

In the United States, where healthcare expenditures reached 18.3% of gross domestic product in 2021 (CMS, 2023), life expectancy remains lower than in numerous countries with substantially smaller investments (Papanicolas et al., 2018). Medical errors constitute the third leading cause of death, with estimates ranging from 44,000 to 250,000 deaths annually depending on methodology and definitions (Makary & Daniel, 2016; James, 2013). The Institute of Medicine's landmark report *To Err Is Human* documented that preventable adverse events affect approximately 3–4% of hospitalizations (Institute of Medicine, 2000), while the subsequent *Crossing the Quality Chasm* report characterized the gap between evidence and practice as a "chasm" rather than a gap (Institute of Medicine, 2001).

In regions with universal healthcare, patients face extended waiting periods, overloaded facilities, and consultations truncated to mere minutes. In low-resource settings, the divide between available knowledge and accessible care becomes pronounced, with populations lacking access to interventions standard for decades elsewhere.

Conventional explanations—inadequate funding, insufficient personnel, bureaucratic inefficiency, educational gaps—while partially valid, fail to address a deeper structural problem. This paper proposes a different diagnosis: *contemporary medicine does not fail because it lacks knowledge; it fails because it cannot structurally carry knowledge through the entire clinical process.*

We introduce the *Clinical Continuity Principle* as a theoretical framework for understanding this fundamental limitation. We argue that effective healthcare requires continuous informational coherence across four interconnected phases: symptom presentation, diagnostic reasoning, therapeutic prescription, and treatment execution. At each transition, human-operated systems experience systematic information loss, responsibility fragmentation, and feedback disruption. These discontinuities are inevitable consequences of cognitive architectures unsuited to the task at scale.

If the core problem is structural rather than circumstantial, solutions targeting funding, training, or protocols yield only marginal improvements. What is required is a mechanism capable of restoring continuity that human cognition cannot maintain. We argue that artificial intelligence represents such a mechanism—not as enhancement, but as the first technology capable of addressing healthcare's architectural deficit. We also acknowledge AI's significant limitations and the substantial barriers to realizing this potential.

1.1. Approach and Methodology

This paper presents a theoretical framework developed through systems analysis and literature synthesis. Our methodology integrates:

(1) **Systems-theoretic analysis** of clinical pathways, identifying transition points and information flow patterns

(2) **Cognitive science literature review** examining human information processing limitations relevant to clinical work

(3) **Empirical synthesis** of published data on error rates, adherence patterns, and information loss in healthcare settings

(4) **Technology assessment** of AI capabilities and limitations in clinical contexts

We do not present original empirical data but synthesize existing evidence to support a novel theoretical framework. The Clinical Continuity Principle is offered as a conceptual contribution requiring future empirical validation across diverse healthcare contexts.

2. The Architecture of Clinical Discontinuity

2.1. Mapping the Clinical Pathway

The clinical pathway from illness to recovery follows phases that should form a coherent continuum but operate as disconnected domains with distinct logics, constraints, and personnel. Figure 1 illustrates this four-phase model with documented discontinuity points.

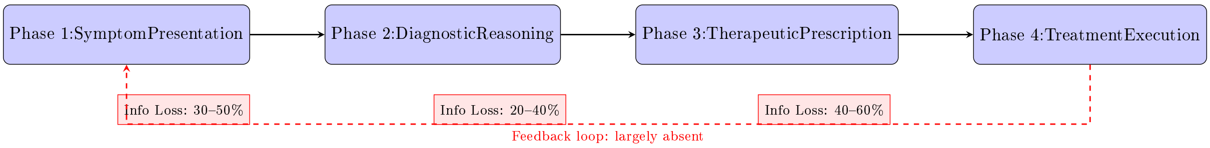

Figure 1: The Four-Phase Clinical Pathway with Documented Discontinuity Points. Information loss estimates represent illustrative ranges synthesized from systematic reviews of diagnostic error, prescription error, and medication adherence literature (Singh et al., 2013; Keers et al., 2013; Brown & Bussell, 2011). These are not results of formal meta-analysis but indicate the scale of documented discontinuities.

The clinical pathway decomposes into four primary phases:

1. **Symptom Presentation**: The patient experiences and communicates bodily changes, discomfort, or concerns. This communication is inherently imprecise, mediated by subjective perception, health literacy, cultural factors, and emotional state.

2. **Diagnostic Reasoning**: The clinician interprets presented information, integrating it with examination findings, laboratory results, and medical history to formulate diagnostic hypotheses.

3. **Therapeutic Prescription**: Based on diagnostic conclusions, the clinician selects interventions—medications, procedures, lifestyle modifications, referrals—intended to address the identified condition.

4. **Treatment Execution**: The patient implements prescribed interventions, monitoring responses and managing complications.

In an idealized system, information would flow seamlessly across phases. Reality diverges dramatically: at each transition, information degrades, responsibilities transfer incompletely, and feedback loops fragment or disappear.

2.2. Empirical Evidence of Discontinuity

Table 1 summarizes empirical evidence for information loss and discontinuity at each clinical transition point.

Table 1: Empirical Evidence for Information Loss at Clinical Transition Points. Note: The aggregate estimates of 30–60% information loss per transition represent illustrative ranges synthesized from the cited studies rather than results of formal meta-analysis. Individual study findings vary by setting, methodology, and population.

Transition	**Finding**	**Magnitude**	**Source**
Symptom → Diagnosis	Diagnostic errors in primary care	5–15% of cases	Singh et al. (2013)
	Missed diagnoses in emergency settings	7.4% of cases	Newman-Toker et al. (2019)
	Relevant history not elicited	30–50% of encounters	Marvel et al. (1999)
Diagnosis → Prescription	Prescribing errors in hospitals	7% of medication orders	Lewis et al. (2009)
	Drug-drug interactions missed	23–35% of potential interactions	Dechanont et al. (2014)
	Guideline non-adherence	30–40% of eligible patients	McGlynn et al. (2003)
Prescription → Execution	Primary non-adherence (never filled)	20–30% of new prescriptions	Fischer et al. (2010)
	Medication non-adherence (chronic)	40–50% of patients	Brown & Bussell (2011)
	Incorrect administration	15–25% of doses	Keers et al. (2013)
Execution → Feedback	Follow-up data not communicated	25–40% of test results	Callen et al. (2012)
	Specialist recommendations lost	20–35% of referrals	Gandhi et al. (2000)

2.3. Case Vignette: Clinical Discontinuity in Practice

Box 1: Illustrative Case of Clinical Discontinuity

Patient: 67-year-old woman with type 2 diabetes, hypertension, and osteoarthritis.
Phase 1 (Symptom Presentation): Reports "feeling tired all the time" and "occasional dizziness" to primary care physician during 12-minute appointment. Does not mention that symptoms worsen after taking morning medications, as she was not specifically asked.
Phase 2 (Diagnostic Reasoning): Physician attributes fatigue to poorly controlled diabetes (HbA1c: 8.2%) and age. Orders routine labs. Does not consider medication side effects or timing-related hypotension. No documentation of symptom timing or aggravating factors.
Phase 3 (Therapeutic Prescription): Physician increases metformin dose to improve glycemic control. Also refills lisinopril, amlodipine, and naproxen (prescribed by rheumatologist). Does not review complete medication list for interactions. Electronic system generates alert for NSAID-ACE inhibitor interaction; physician clicks through without action due to "alert fatigue."
Phase 4 (Treatment Execution): Patient fills metformin prescription but discontinues after one week due to gastrointestinal side effects, which she attributes to "the diabetes getting worse." Does not contact physician. Continues taking naproxen with morning antihypertensives, contributing to postural hypotension (the actual cause of dizziness).
Feedback Loop: Three months later, patient presents to emergency department after fall with hip fracture. Only at this point is the medication interaction identified. No mechanism existed to detect the intervening treatment failure.
Discontinuity Analysis: Information lost at each transition; no integrated oversight of medication regimen; patient unable to navigate system complexity; feedback arrived only after serious harm.

2.4. Transition One: From Symptom to Diagnosis

The first discontinuity occurs at the interface between patient experience and clinical interpretation. Studies demonstrate that physicians interrupt patients within 11–18 seconds of opening statements (Marvel et al., 1999), and that up to 54% of patient complaints and 45% of patient concerns remain undisclosed during encounters (Barry et al., 2001).

The diagnostic process operates under significant constraints. Cognitive psychology has documented numerous heuristics and biases affecting clinical reasoning (Croskerry, 2009): anchoring on initial impressions, premature closure terminating hypothesis generation too early, availability bias overweighting recently encountered conditions, and confirmation bias selectively attending to evidence supporting favored diagnoses. A systematic review by Saposnik et al. (2016) found that cognitive biases contributed to diagnostic error in 36.5–77% of cases examined.

Moreover, medical knowledge volume has rendered comprehensive consideration impossible for any individual clinician. Densen (2011) estimated that maintaining currency would require reading 29 hours daily. Even conscientious practitioners operate within bounded knowledge subsets shaped by training, specialty, and recent experience.

2.5. Transition Two: From Diagnosis to Prescription

The second discontinuity manifests between diagnostic conclusion and therapeutic selection. McGlynn et al. (2003) found that patients received only 54.9% of recommended care across 439 quality indicators, with adherence ranging from 78.7% (senile cataract) to 10.5% (alcohol dependence).

Drug interactions represent a consequential aspect of this discontinuity. As polypharmacy becomes increasingly common—approximately 40% of adults over 65 take five or more medications (Kantor et al., 2015)—combinatorial complexity exceeds human computational capacity. Dechanont et al. (2014) found that drug-drug interactions caused 0.054–0.57% of emergency department visits and 0.57–4.76% of hospital admissions, many preventable with systematic checking.

Furthermore, prescribing clinicians typically lack visibility into previous treatment attempts, responses, or the patient's complete medication landscape across multiple prescribers. Gandhi et al. (2000) documented that 25% of specialist recommendations were never implemented by referring physicians.

2.6. Transition Three: From Prescription to Execution

The third discontinuity—perhaps most consequential yet least examined—occurs between prescription and treatment implementation. Once a prescription is written, clinician involvement typically ceases. Fischer et al. (2010) found that 20–30% of new prescriptions are never filled ("primary non-adherence"), while Brown & Bussell (2011) documented that adherence rates for chronic disease treatments average only 50%.

Patients discontinue medications due to side effects, cost, complexity, or skepticism. They modify dosages based on perceived response or non-medical advice. The World Health Organization characterized poor adherence as "a worldwide problem of striking magnitude," estimating that interventions addressing adherence would have greater population health impact than improvements in specific treatments (WHO, 2003).

The prescribing physician receives no notification of these deviations. The feedback loop connecting treatment response to therapeutic adjustment is fundamentally broken. Clinicians learn of treatment failures only when patients return—often after significant deterioration—or sometimes never.

2.7. Information Entropy in Clinical Systems

The cumulative effect of these discontinuities can be understood through information theory. At each transition, the clinical process experiences information entropy—meaningful data degradation. Initial symptom presentations contain rich, if unstructured, information. Through successive translations and handoffs, this information becomes compressed, filtered, and corrupted.

By the time treatment reaches execution, connection to original clinical reality becomes tenuous. The medication prescribed reflects a diagnosis that simplified symptom complexity, based on guidelines that generalized individual variation, without consideration of factors determining actual adherence and response. The gap between theoretical treatment and realized therapeutic effect represents accumulated entropy of an intrinsically discontinuous system.

3. Cognitive Constraints and Structural Necessity

3.1. The Limits of Human Information Processing

The clinical discontinuities described are not primarily failures of individual competence, institutional design, or resource allocation. They are inevitable consequences of maintaining informational coherence through human cognitive systems operating at modern healthcare's scale and complexity.

Table 2 summarizes relevant cognitive limitations and their clinical manifestations.

Table 2: Human Cognitive Limitations and Clinical Manifestations

Cognitive Limitation	**Description**	**Clinical Manifestation**	**Source**
Working memory capacity	Approximately 4 chunks maintained simultaneously (Cowan's model)	Cannot hold complete patient context during decision-making	Cowan (2010)
Attentional limitations	Selective attention; dual-task interference	Documentation competes with patient observation	Pashler (1994)
Long-term memory decay	Encoding requires effort; retrieval degrades	Patient details forgotten between encounters	Ebbinghaus (1885)
Processing speed	50 bits/second conscious processing	Cannot evaluate all drug interactions in real-time	Miller (1956)
Decision fatigue	Quality degrades with successive decisions	Later patients receive less thorough evaluation	Linder et al. (2014)
Cognitive biases	Systematic deviations from optimal reasoning	Diagnostic anchoring, premature closure	Croskerry (2009)

According to Cowan (2010)'s influential model, human working memory can maintain approximately four chunks of information simultaneously. Long-term memory, while capacious, requires encoding effort and degrades over time. Attention is inherently selective, unable to simultaneously monitor multiple information streams. Processing speed limits the rate at which complex information can be analyzed.

Consider the cognitive demands facing a clinician in routine practice. During a typical consultation, the physician must simultaneously attend to verbal communication, observe nonverbal cues, formulate questions, recall pertinent knowledge, integrate current findings with historical information, consider differential diagnoses, evaluate treatment options, document the encounter, and manage time constraints. Sinsky et al. (2016) found that for every hour of direct clinical face time, physicians spend nearly two additional hours on electronic health record tasks.

Linder et al. (2014) demonstrated decision fatigue in clinical settings: antibiotic prescribing for acute respiratory infections increased throughout clinical sessions, suggesting that decision quality degrades with successive choices. Multiply this burden across dozens of patients daily, and the accumulation creates impossible demands on human memory and attention.

3.2. The Myth of Compensatory Systems

Healthcare systems have developed mechanisms intended to compensate for individual cognitive limitations: protocols, checklists, electronic health records, clinical decision support systems, team-based care models. While these interventions reduce specific error types, they cannot eliminate the fundamental discontinuity problem.

Electronic health records promised to solve information discontinuity by creating comprehensive, accessible documentation. In practice, EHR systems often impede clinical work. Ratwani et al. (2018) documented usability problems contributing to safety events. Campbell et al. (2006) identified nine categories of unintended EHR consequences, including more/new work, workflow issues, and communication problems. Physicians report spending more time with computers than patients, without corresponding information quality improvements (Sinsky et al., 2016).

Clinical decision support systems generate thousands of alerts, but "alert fatigue" leads clinicians to override 49–96% of drug interaction warnings (van der Sijs et al., 2006). The cognitive load of attending to multiple reminder systems itself becomes a distraction from patient care.

These compensatory mechanisms fail not because of poor design but because they cannot transcend the fundamental constraint: maintaining clinical continuity requires cognitive capacities that humans do not possess. The problem is architectural, not operational.

3.3. The Clinical Continuity Principle

We propose the *Clinical Continuity Principle* as the foundational requirement for effective healthcare delivery:

> *Effective medical care requires continuous informational coherence across all stages of the clinical process—from initial symptom presentation through diagnostic reasoning, therapeutic prescription, and treatment execution. This coherence must be maintained at the level of individual patients while simultaneously integrating population-level evidence and system-wide coordination.*

This principle has several implications:

- **Comprehensiveness**: All relevant information must be accessible and considered at each decision point.
- **Integration**: Information from different sources and time points must be synthesized into coherent understanding.
- **Persistence**: The clinical process extends beyond individual encounters; treatment effects must be continuously monitored.
- **Feedback**: Outcomes must flow back to inform future decisions.

Human cognitive architecture cannot satisfy these requirements at scale. The working memory limitations, attentional constraints, and processing bottlenecks characterizing human cognition make comprehensive, integrated, persistent, feedback-rich clinical continuity structurally impossible when care is mediated entirely through human agents.

4. Artificial Intelligence as Continuity Mechanism

4.1. Beyond Optimization: AI as Architectural Solution

The discourse surrounding AI in healthcare has predominantly framed AI as an optimization tool—making existing processes faster or more accurate. Diagnostic algorithms improve radiology interpretation (McKinney et al., 2020); predictive models identify high-risk patients (Rajkomar et al., 2018); natural language processing extracts information from clinical notes (Jiang et al., 2017). These applications conceptualize AI as enhancement to fundamentally human-centered systems.

The analysis presented here suggests a more fundamental role. If modern medicine's core problem is structural discontinuity arising from human cognitive limitations, then AI represents not optimization but architectural transformation—the first mechanism capable of maintaining clinical continuity that human systems cannot achieve.

This reframing shifts attention from incremental improvements within existing frameworks to fundamental reconceptualization of healthcare organization.

4.2. Characteristics Enabling Continuity

Table 3 contrasts human cognitive limitations with AI capabilities relevant to clinical continuity.

Table 3: Comparison of Human Cognitive Limitations and AI Capabilities for Clinical Continuity. Note: AI capabilities listed represent theoretical potential; actual performance depends on system design, training data quality, and implementation context. AI systems have their own limitations including potential for hallucination, encoded biases, and brittleness under distribution shift.

Dimension	Human Limitation	AI Capability (Theoretical)
Working memory	$\sim$4 chunks simultaneously	Large active context; complete patient information maintainable (limited by context window)
Knowledge access	Cannot retain complete medical literature; bounded by training and experience	Can access comprehensive databases; but may hallucinate or retrieve outdated information
Attention consistency	Fatigue, distraction, attentional fluctuation; quality degrades through shift	Consistent processing regardless of volume; no fatigue (but may have systematic blind spots)
Monitoring persistence	Episodic attention; patients "invisible" between encounters	Continuous observation possible; real-time trend analysis (requires infrastructure)
Recall accuracy	Memory degrades; details lost over time	Perfect digital recall of recorded data; no degradation
Processing scale	Limited to sequential patient encounters	Parallel analysis of many cases; population-level pattern recognition
Bias susceptibility	Systematic cognitive biases affect reasoning	Different bias profile: may encode training data biases; can be designed to mitigate specific human biases

AI systems possess characteristics directly addressing identified discontinuity problems:

Unlimited working memory: AI systems can maintain active representations of all patient-relevant information simultaneously. Complete medical history, current symptoms, active medications, pending test results, relevant guidelines, and population-level evidence can inform each decision without trade-offs imposed by human capacity limits.

Comprehensive knowledge access: Machine learning systems can be trained on complete medical literature, clinical trial data, and accumulated case records. Knowledge that no individual physician could master becomes accessible to every patient interaction.

Consistent attention: Unlike human cognition, AI systems do not experience fatigue, distraction, or attentional fluctuation. The fortieth patient receives the same comprehensive analysis as the first.

Persistent monitoring: AI systems can maintain continuous observation of patient status, tracking biomarker trends, adherence patterns, and symptom evolution without the intermittency of episodic human attention.

Perfect recall: Every prior interaction, recorded value, and documented decision remains accessible without degradation.

Scalable processing: The same analytical capabilities apply to unlimited simultaneous patients without degradation.

4.3. Evidence for AI Clinical Performance

Emerging evidence supports AI capabilities in specific clinical domains:

- **Diagnostic imaging**: McKinney et al. (2020) demonstrated AI performance exceeding average radiologist accuracy in breast cancer screening, with absolute reductions in false positives (5.7%) and false negatives (9.4%).

- **Diagnostic reasoning**: Liu et al. (2019) showed deep learning systems achieving diagnostic accuracy comparable to dermatologists for skin lesion classification.
- **Risk prediction**: Rajkomar et al. (2018) demonstrated that deep learning models using EHR data predicted in-hospital mortality, 30-day unplanned readmission, and prolonged length of stay with AUCs of 0.93–0.95.
- **Medication management**: AI-based clinical decision support has demonstrated reductions in adverse drug events when properly implemented (Roshanov et al., 2011).

However, we note that most evidence comes from controlled settings, and real-world implementation has produced mixed results (see Section 5).

4.4. Restoring Continuity Across Transitions

The structural discontinuities identified earlier can be specifically addressed through AI capabilities:

Symptom to Diagnosis: Natural language processing systems can interpret patient descriptions, identifying clinically relevant features even when expressed in non-medical vocabulary. Pattern recognition can surface diagnostic possibilities that might not occur to attending physicians, particularly rare conditions or atypical presentations.

Diagnosis to Prescription: AI systems can evaluate treatment options against complete patient profiles, including all comorbidities, current medications, genetic factors, and prior treatment responses. Drug interaction checking can be comprehensive rather than superficial.

Prescription to Execution: Connected monitoring systems can track adherence, detecting when medications are not taken as prescribed. Automated follow-up can identify side effects before discontinuation. Feedback from wearable devices and patient-reported outcomes can be continuously integrated.

Execution to Adjustment: The closed loop from treatment execution back to diagnostic reassessment and therapeutic modification—rarely achieved in human-operated systems—becomes structurally possible.

4.5. The Physician-AI Partnership

Restoring clinical continuity through AI does not imply replacing physicians. Human clinicians retain essential roles: interpersonal connection supporting patient trust and adherence, ethical judgment for complex decisions, creative flexibility for unprecedented situations, and communicative skills translating clinical understanding into patient comprehension.

What changes is clinical work's cognitive architecture. Rather than physicians attempting to maintain continuity through intrinsically limited cognitive resources, they operate within an environment where continuity is structurally assured. The AI system maintains comprehensive context, monitors progress, identifies relevant evidence, and flags developing problems. The physician focuses on essentially human elements: communicating with patients, making value-laden decisions, providing interpersonal care.

This partnership addresses concerns about depersonalization. Paradoxically, when physicians are relieved of the impossible cognitive burden of maintaining clinical continuity, they gain capacity for genuinely personal aspects of medicine. Attention previously divided between patient interaction and information management can concentrate on the patient as a person.

5. Limitations, Counterarguments, and Barriers

5.1. AI Reliability and Error

AI systems are not infallible. Significant concerns include:

Hallucination and confabulation: Large language models can generate plausible-sounding but factually incorrect information (Ji et al., 2023). In clinical contexts, this could produce dangerous recommendations.

Distribution shift: AI systems trained on historical data may perform poorly when clinical patterns change, as demonstrated during COVID-19 when many prediction models failed in deployment (Wynants et al., 2020).

Brittleness: Small input perturbations can produce dramatically different outputs, and adversarial examples can systematically mislead AI systems (Finlayson et al., 2019).

Failure mode opacity: When AI systems err, identifying why may be difficult, complicating quality improvement.

We acknowledge that AI does not eliminate error but potentially transforms its character. Human errors are often random and idiosyncratic; AI errors may be systematic and correlated across patients. Appropriate safeguards require understanding these different error profiles.

5.2. Algorithmic Bias and Equity

AI systems can encode and amplify biases present in training data:

Historical bias: If training data reflects historical disparities in care, AI may perpetuate them. Obermeyer et al. (2019) demonstrated that a widely-used algorithm for identifying high-risk patients systematically underestimated illness severity for Black patients because it used healthcare costs as a proxy for health needs.

Representation bias: If training data underrepresents certain populations, AI performance may degrade for those groups.

Measurement bias: If outcome measures are systematically different across groups (e.g., due to differential access to diagnosis), learned associations may be misleading.

Addressing algorithmic bias requires active intervention: diverse training data, fairness-aware learning algorithms, ongoing monitoring for disparate impact, and diverse development teams. The potential for AI to democratize access to high-quality care is real but will not be realized automatically.

5.3. The "Black Box" Problem

Many high-performing AI systems, particularly deep learning models, operate as "black boxes" whose reasoning cannot be easily explained. This creates challenges:

Clinical trust: Physicians may be reluctant to follow recommendations they cannot understand or verify.

Patient autonomy: Patients cannot meaningfully consent to AI-influenced care if the basis for recommendations is opaque.

Accountability: When AI contributes to adverse outcomes, determining responsibility is complicated by inability to audit reasoning.

Quality improvement: Learning from errors requires understanding why they occurred.

Monitoring fatigue: Paradoxically, the need to critically evaluate AI recommendations may create new cognitive burdens, replacing one form of overload with another. Physicians may face "AI oversight fatigue" analogous to alert fatigue in current clinical decision support systems.

Explainable AI (XAI) research addresses some of these concerns (Amann et al., 2020), and emerging techniques such as retrieval-augmented generation (RAG) can ground AI outputs in verifiable sources. However, trade-offs between performance and interpretability remain. We suggest that different clinical contexts may warrant different positions on this trade-off: screening applications may tolerate less explainability than treatment recommendations. Human-in-the-loop oversight remains essential for high-stakes decisions.

5.4. The Hybrid Human-AI Reality

It is important to acknowledge that pure AI-maintained continuity is neither realistic nor desirable in the near term. The practical path forward involves hybrid human-AI systems in which AI augments rather than replaces human oversight. Emerging technologies such as ambient AI scribes, predictive monitoring dashboards, and intelligent care coordination platforms already bridge some continuity gaps without full automation.

This hybrid model introduces its own challenges. If AI addresses some discontinuities while introducing new structural problems—workflow integration failures, data fragmentation across AI tools, and the cognitive burden of monitoring AI outputs—continuity may be shifted rather than restored. Success requires designing AI systems that genuinely reduce net cognitive load while improving information flow, rather than simply adding another layer of complexity to already overburdened clinical environments.

5.5. Implementation Barriers

Translating AI capabilities into clinical practice faces substantial barriers:

Integration with workflows: AI systems must fit into clinical workflows without creating additional burden. Many promising technologies have failed at this stage (Kelly et al., 2019).

Interoperability: Healthcare data exists in fragmented, incompatible systems. Realizing AI continuity benefits requires data integration that current infrastructure often precludes.

Regulatory uncertainty: Regulatory frameworks for AI in healthcare remain underdeveloped, creating uncertainty for developers and implementers.

Economic models: Fee-for-service reimbursement may not support AI-enabled continuous care. Payment reform may be prerequisite to realizing continuity benefits.

Workforce adaptation: Clinicians require training to work effectively with AI systems. Resistance to practice change is well-documented.

Digital divide: AI-dependent care could disadvantage patients lacking technology access or digital literacy, potentially widening rather than narrowing health disparities.

5.6. Real-World Implementation Experience

Early implementations have produced mixed results:

IBM Watson for Oncology: Despite substantial investment and initial enthusiasm, Watson for Oncology faced criticism for providing recommendations inconsistent with expert consensus and has been discontinued by many adopting institutions (Ross & Swetlitz, 2018).

Epic Sepsis Model: An analysis by Wong et al. (2021) found that Epic's widely-deployed sepsis prediction model performed substantially worse than reported in vendor materials when evaluated externally.

Positive implementations: Conversely, AI-assisted diabetic retinopathy screening has demonstrated effective real-world deployment in multiple settings (Beede et al., 2020), and some clinical decision support systems have demonstrated sustained effectiveness (Roshanov et al., 2011).

These mixed results underscore that AI capability does not automatically translate to clinical benefit. Successful implementation requires attention to workflow integration, user training, appropriate use cases, and ongoing monitoring.

6. Implications and Implementation Considerations

6.1. Healthcare System Transformation

The framework presented implies that healthcare systems should be reconceptualized around AI-enabled continuity capabilities. Rather than integrating AI tools into existing workflows, workflows should be redesigned to leverage AI-maintained continuity as foundational infrastructure.

This transformation involves rethinking traditional boundaries. The distinction between inpatient and outpatient care, between acute treatment and chronic management, between different specialty domains—all reflect organizational conventions that fragment clinical continuity. AI systems can maintain integrated patient representations across these artificial boundaries.

Similarly, the episodic nature of clinical encounters—consultations scheduled at administrative convenience—gives way to continuous monitoring punctuated by human interaction when genuinely needed.

6.2. Equity Considerations

The clinical continuity problem disproportionately affects vulnerable populations. Patients with limited health literacy struggle to navigate fragmented systems. Those with multiple chronic conditions face severe coordination challenges. Residents of underserved areas lack specialist access.

AI-maintained clinical continuity has potential to reduce these disparities. The same comprehensive analytical capabilities can be deployed regardless of patient socioeconomic status or geographic location. However, this equity potential will not be realized automatically. Active policy decisions are required to ensure AI applications are deployed with equity as priority rather than reinforcing existing advantages.

The digital divide represents a particular concern. Patients lacking smartphone access, internet connectivity, or digital literacy could be disadvantaged by AI-dependent care models. Implementation strategies must include provisions for technology access and non-digital alternatives.

6.3. Regulatory and Ethical Frameworks

The transformation envisioned requires evolution in regulatory frameworks. Current structures assume human clinicians making decisions based on individual judgment. Accountability, liability, and quality assurance mechanisms assume human agents at each decision point.

Integrating AI raises questions current frameworks do not adequately address:

- When AI recommendations influence clinical decisions, how is responsibility allocated?
- What transparency requirements should apply to algorithmic reasoning?
- How should AI systems be validated and monitored for safety?
- How is patient autonomy protected when AI systems access comprehensive personal data?

Privacy protections represent a particular challenge that merits detailed consideration. Clinical continuity benefits depend on comprehensive data access—the more complete the patient picture, the more effectively AI can maintain coherence across transitions. Yet comprehensive access creates significant privacy risks: data breaches, unauthorized secondary uses, surveillance concerns, and potential discrimination based on health information.

This tension between informational continuity and privacy is not merely technical but fundamentally ethical. The Clinical Continuity Principle, taken to its logical extreme, would favor maximizing data integration. But patient autonomy and dignity require meaningful control over personal health information. Different patients may reasonably reach different conclusions about this trade-off based on their values, circumstances, and risk tolerance.

Technical solutions—differential privacy, federated learning, secure multi-party computation, and data minimization principles—can mitigate some risks while preserving analytical capabilities. Governance structures including patient consent frameworks, data access auditing, and independent oversight can provide accountability. But fundamental tensions between informational continuity and privacy require ongoing societal negotiation rather than purely technical resolution. The framework we propose must be implemented within bounds that respect patient autonomy and democratic decision-making about acceptable data practices.

6.4. Implementation Pathways

The transition to AI-maintained clinical continuity cannot occur instantaneously. Pragmatic implementation requires identifying entry points where AI continuity capabilities can demonstrate value while building toward more comprehensive transformation.

Chronic disease management: Patients with conditions like diabetes, heart failure, or COPD require continuous monitoring and adjustment that AI systems can provide. Current chronic care fragmentation exemplifies the discontinuity problem clearly.

Medication management: The prescription-to-execution discontinuity generates substantial morbidity through adverse drug events, non-adherence, and drug interactions. AI systems maintaining comprehensive medication profiles can address these problems while demonstrating broader capabilities.

Care transitions: Hospital discharge transitions represent high-risk discontinuity points where AI continuity management could prevent readmissions that current fragmented handoffs cause.

Surgical pathways: From preoperative evaluation through postoperative recovery, surgical care benefits from continuous monitoring and coordination that AI enables.

7. The Epistemological Dimension

7.1. Beyond Application: AI as Discovery Engine

The discussion thus far has focused on AI as a mechanism for applying existing medical knowledge more effectively. But implications extend beyond application to discovery. The same capabilities that enable clinical continuity—comprehensive data integration, pattern recognition across vast datasets, identification of subtle correlations—position AI as a tool for generating new medical knowledge.

Diseases currently treated as unified entities may prove to be heterogeneous collections of distinct processes requiring different interventions. AI systems analyzing large populations can identify subtypes, distinguishing patients who appear clinically similar but respond differently to treatment. Conversely, conditions currently classified separately may share underlying mechanisms that AI analysis reveals.

Treatment responses that appear idiosyncratic may reflect unrecognized patterns. Why does a medication that helps most patients fail dramatically for some? The answer often lies in combinations of factors—genetic variants, microbiome composition, lifestyle factors, comorbidities—too complex for human analysis but tractable for AI systems integrating diverse data sources.

7.2. Transforming Medical Education

If clinical practice transforms as described, medical education must transform correspondingly. Current training emphasizes developing individual clinical reasoning capabilities. These capabilities remain valuable but no longer suffice for healthcare systems in which AI maintains clinical continuity.

Future clinicians must learn to work within human-AI collaborative frameworks: understanding AI capabilities and limitations, interpreting AI-generated recommendations critically, identifying situations requiring human judgment that AI cannot provide, and communicating effectively with patients about AI roles in their care.

7.3. Testable Hypotheses and Future Research

The Clinical Continuity Principle, as a theoretical framework, generates testable hypotheses that future empirical research should address:

1. **Continuity Index Hypothesis**: A quantifiable "Clinical Continuity Index" measuring information retention across pathway transitions should correlate with patient outcomes. Research could develop and validate such metrics, comparing AI-augmented versus traditional care pathways.

2. **Transition-Specific Intervention Hypothesis**: Interventions targeting specific discontinuity points (e.g., AI-mediated prescription-to-execution monitoring) should produce measurable improvements in that transition's information retention, with downstream effects on outcomes.

3. **Cognitive Load Trade-off Hypothesis**: AI systems designed for continuity maintenance may reduce certain cognitive burdens while introducing others (e.g., monitoring fatigue). Net cognitive load and its relationship to error rates should be empirically assessed.

4. **Equity Impact Hypothesis**: AI-maintained continuity should disproportionately benefit populations currently most harmed by fragmentation (e.g., patients with multiple chronic conditions, limited health literacy, or in underserved areas), assuming equitable access to the technology.

5. **Feedback Loop Restoration Hypothesis**: AI systems that close the execution-to-feedback loop through continuous monitoring should demonstrate improved treatment adjustment rates and reduced time-to-intervention compared to episodic care models.

These hypotheses provide a research agenda for moving beyond theoretical argument to empirical validation. Randomized controlled trials comparing AI-augmented continuity interventions against standard care, using validated continuity metrics as intermediate outcomes and patient health outcomes as primary endpoints, would provide the strongest evidence for or against the framework's core claims.

8. Conclusion

Contemporary healthcare suffers from a fundamental architectural deficit. The clinical process—from symptom presentation through diagnosis, prescription, and treatment execution—requires continuous informational coherence that human cognitive systems cannot maintain at scale. Knowledge adequate to transform patient outcomes exists but cannot flow reliably through structural discontinuities characterizing human-operated healthcare.

Artificial intelligence emerges from this analysis not as one technology among many but as a mechanism with unique potential for restoring clinical continuity. The comprehensive memory, consistent attention, persistent monitoring, and integrated analysis that AI systems can provide address precisely the cognitive limitations creating discontinuity in human-mediated care. However, we emphasize that the realistic path forward involves hybrid human-AI systems rather than full automation, and that AI introduces its own challenges including potential for hallucination, encoded biases, and the risk of creating new cognitive burdens for clinicians.

We have been careful to acknowledge AI's significant limitations: reliability concerns, algorithmic bias, the black box problem, implementation barriers, and the fundamental tension between informational continuity and patient privacy. Mixed real-world results underscore that capability does not automatically translate to benefit. The transformation we describe is not inevitable but contingent on addressing these challenges through careful system design, appropriate governance, and ongoing empirical evaluation.

The Clinical Continuity Principle we propose is offered as a theoretical framework requiring empirical validation. We have articulated testable hypotheses that future research should address, including development of quantifiable continuity metrics and randomized trials comparing AI-augmented versus traditional care pathways. This paper is a conceptual provocation rather than a final answer.

Nevertheless, the direction indicated by this analysis is clear. Medicine does not fail because it lacks knowledge. It fails because it cannot structurally carry knowledge through the entire clinical process. Artificial intelligence, thoughtfully implemented within hybrid human-AI systems, offers the possibility of closing that structural gap—not through replacement of human judgment and care, but through partnership that combines AI's continuity capabilities with human clinical wisdom and interpersonal connection.

The era beginning is one in which health becomes not merely an outcome pursued through episodic interventions but a state continuously maintained through intelligent, integrated systems. In this transformation lies the prospect of healthcare that finally delivers on the promise medical knowledge has long implied but fragmented delivery systems have chronically frustrated.

Declarations

Funding

This work received no external funding.

Conflicts of Interest

The author declares no conflicts of interest.

Data Availability

This is a theoretical framework paper. No original data were collected. All referenced data are from published sources cited in the text.

Author Contributions

B.K. conceived the theoretical framework, conducted the literature synthesis, and wrote the manuscript.

Ethics Approval

Not applicable. This paper presents a theoretical framework based on literature synthesis and does not involve human subjects research.

References

Amann, J., Blasimme, A., Vayena, E., Frey, D., & Madai, V. I. (2020). Explainability for artificial intelligence in healthcare: A multidisciplinary perspective. *BMC Medical Informatics and Decision Making*, 20(1), 310.

Barry, C. A., Bradley, C. P., Britten, N., Stevenson, F. A., & Barber, N. (2001). Patients' unvoiced agendas in general practice consultations: Qualitative study. *BMJ*, 320(7244), 1246–1250.

Beede, E., Baylor, E., Hersch, F., et al. (2020). A human-centered evaluation of a deep learning system deployed in clinics for the detection of diabetic retinopathy. *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, 1–12.

Brown, M. T., & Bussell, J. K. (2011). Medication adherence: WHO cares? *Mayo Clinic Proceedings*, 86(4), 304–314.

Callen, J. L., Westbrook, J. I., Georgiou, A., & Li, J. (2012). Failure to follow-up test results for ambulatory patients: A systematic review. *Journal of General Internal Medicine*, 27(10), 1334–1348.

Campbell, E. M., Sittig, D. F., Ash, J. S., Guappone, K. P., & Dykstra, R. H. (2006). Types of unintended consequences related to computerized provider order entry. *Journal of the American Medical Informatics Association*, 13(5), 547–556.

Centers for Medicare & Medicaid Services. (2023). National health expenditure data: Historical. Retrieved from https://www.cms.gov/research-statistics-data-and-systems/statistics-trends-and-reports/nationalhealthexpenddata

Cowan, N. (2010). The magical mystery four: How is working memory capacity limited, and why? *Current Directions in Psychological Science*, 19(1), 51–57.

Croskerry, P. (2009). A universal model of diagnostic reasoning. *Academic Medicine*, 84(8), 1022–1028.

Dechanont, S., Maphanta, S., Butthum, B., & Kongkaew, C. (2014). Hospital admissions/visits associated with drug-drug interactions: A systematic review and meta-analysis. *Pharmacoepidemiology and Drug Safety*, 23(5), 489–497.

Densen, P. (2011). Challenges and opportunities facing medical education. *Transactions of the American Clinical and Climatological Association*, 122, 48–58.

Ebbinghaus, H. (1885). *Über das Gedächtnis: Untersuchungen zur experimentellen Psychologie.* Leipzig: Duncker & Humblot.

Finlayson, S. G., Bowers, J. D., Ito, J., Zittrain, J. L., Beam, A. L., & Kohane, I. S. (2019). Adversarial attacks on medical machine learning. *Science*, 363(6433), 1287–1289.

Fischer, M. A., Stedman, M. R., Lii, J., et al. (2010). Primary medication non-adherence: Analysis of 195,930 electronic prescriptions. *Journal of General Internal Medicine*, 25(4), 284–290.

Gandhi, T. K., Sittig, D. F., Franklin, M., Sussman, A. J., Fairchild, D. G., & Bates, D. W. (2000). Communication breakdown in the outpatient referral process. *Journal of General Internal Medicine*, 15(9), 626–631.

Institute of Medicine. (2000). *To Err Is Human: Building a Safer Health System.* Washington, DC: National Academies Press.

Institute of Medicine. (2001). *Crossing the Quality Chasm: A New Health System for the 21st Century.* Washington, DC: National Academies Press.

James, J. T. (2013). A new, evidence-based estimate of patient harms associated with hospital care. *Journal of Patient Safety*, 9(3), 122–128.

Ji, Z., Lee, N., Frieske, R., et al. (2023). Survey of hallucination in natural language generation. *ACM Computing Surveys*, 55(12), 1–38.

Jiang, F., Jiang, Y., Zhi, H., et al. (2017). Artificial intelligence in healthcare: Past, present and future. *Stroke and Vascular Neurology*, 2(4), 230–243.

Kantor, E. D., Rehm, C. D., Haas, J. S., Chan, A. T., & Giovannucci, E. L. (2015). Trends in prescription drug use among adults in the United States from 1999–2012. *JAMA*, 314(17), 1818–1831.

Keers, R. N., Williams, S. D., Cooke, J., & Ashcroft, D. M. (2013). Causes of medication administration errors in hospitals: A systematic review of quantitative and qualitative evidence. *Drug Safety*, 36(11), 1045–1067.

Kelly, C. J., Karthikesalingam, A., Suleyman, M., Corrado, G., & King, D. (2019). Key challenges for delivering clinical impact with artificial intelligence. *BMC Medicine*, 17(1), 195.

Lewis, P. J., Dornan, T., Taylor, D., Tully, M. P., Wass, V., & Ashcroft, D. M. (2009). Prevalence, incidence and nature of prescribing errors in hospital inpatients: A systematic review. *Drug Safety*, 32(5), 379–389.

Linder, J. A., Doctor, J. N., Friedberg, M. W., et al. (2014). Time of day and the decision to prescribe antibiotics. *JAMA Internal Medicine*, 174(12), 2029–2031.

Liu, Y., Jain, A., Eng, C., et al. (2019). A deep learning system for differential diagnosis of skin diseases. *Nature Medicine*, 26(6), 900–908.

Makary, M. A., & Daniel, M. (2016). Medical error—The third leading cause of death in the US. *BMJ*, 353, i2139.

Marvel, M. K., Epstein, R. M., Flowers, K., & Beckman, H. B. (1999). Soliciting the patient's agenda: Have we improved? *JAMA*, 281(3), 283–287.

McGlynn, E. A., Asch, S. M., Adams, J., et al. (2003). The quality of health care delivered to adults in the United States. *New England Journal of Medicine*, 348(26), 2635–2645.

McKinney, S. M., Sieniek, M., Godbole, V., et al. (2020). International evaluation of an AI system for breast cancer screening. *Nature*, 577(7788), 89–94.

Miller, G. A. (1956). The magical number seven, plus or minus two: Some limits on our capacity for processing information. *Psychological Review*, 63(2), 81–97.

Newman-Toker, D. E., Schaffer, A. C., Yu-Moe, C. W., et al. (2019). Serious misdiagnosis-related harms in malpractice claims: The "Big Three"—vascular events, infections, and cancers. *Diagnosis*, 6(3), 227–240.

Obermeyer, Z., Powers, B., Vogeli, C., & Mullainathan, S. (2019). Dissecting racial bias in an algorithm used to manage the health of populations. *Science*, 366(6464), 447–453.

Papanicolas, I., Woskie, L. R., & Jha, A. K. (2018). Health care spending in the United States and other high-income countries. *JAMA*, 319(10), 1024–1039.

Pashler, H. (1994). Dual-task interference in simple tasks: Data and theory. *Psychological Bulletin*, 116(2), 220–244.

Rajkomar, A., Oren, E., Chen, K., et al. (2018). Scalable and accurate deep learning with electronic health records. *NPJ Digital Medicine*, 1(1), 18.

Ratwani, R. M., Reber, E., Grayson, M., & Sheehan, B. (2018). A usability and safety analysis of electronic health records: A multi-center study. *Journal of the American Medical Informatics Association*, 25(9), 1197–1201.

Roshanov, P. S., Misra, S., Gerber, H. M., et al. (2011). Computerized clinical decision support systems for chronic disease management: A decision-maker-researcher partnership systematic review. *Implementation Science*, 6(1), 92.

Ross, C., & Swetlitz, I. (2018). IBM's Watson supercomputer recommended "unsafe and incorrect" cancer treatments, internal documents show. *STAT News*, July 25, 2018.

Saposnik, G., Redelmeier, D., Ruff, C. C., & Tobler, P. N. (2016). Cognitive biases associated with medical decisions: A systematic review. *BMC Medical Informatics and Decision Making*, 16(1), 138.

Singh, H., Giardina, T. D., Meyer, A. N., Forjuoh, S. N., Reis, M. D., & Thomas, E. J. (2013). Types and origins of diagnostic errors in primary care settings. *JAMA Internal Medicine*, 173(6), 418–425.

Sinsky, C., Colligan, L., Li, L., et al. (2016). Allocation of physician time in ambulatory practice: A time and motion study in 4 specialties. *Annals of Internal Medicine*, 165(11), 753–760.

van der Sijs, H., Aarts, J., Vulto, A., & Berg, M. (2006). Overriding of drug safety alerts in computerized physician order entry. *Journal of the American Medical Informatics Association*, 13(2), 138–147.

World Health Organization. (2003). *Adherence to Long-Term Therapies: Evidence for Action*. Geneva: World Health Organization.

Wong, A., Otles, E., Donnelly, J. P., et al. (2021). External validation of a widely implemented proprietary sepsis prediction model in hospitalized patients. *JAMA Internal Medicine*, 181(8), 1065–1070.

Wynants, L., Van Calster, B., Collins, G. S., et al. (2020). Prediction models for diagnosis and prognosis of COVID-19: Systematic review and critical appraisal. *BMJ*, 369, m1328.